

Idea Matrix

THCMCP

The debate/verification engine and its coding pillar — status view as of 2026-09-16.

67 ideas tracked. 22 built · 5 partial · 13 planned · 23 idea · 2 deferred · 2 declined

Legend: ■ BUILT — shipped · ■ PARTIAL — built, gaps remain · ■ PLANNED — scoped, not built · ■ IDEA — not scoped
· ■ DEFERRED — has a named reopen trigger · ■ DECLINED — real reason on file

Debate & verification mechanism

Idea	Status	Notes	Source
Tool-grounded verification	■ BUILT	Run tests/checks, inject ground truth into panel review. The single highest-value gap the research found — closed.	v7 core
Evaluation harness vs. single strong model	■ PARTIAL	Gate item ranked above all v7 features. SWE-bench diff-review pilot planned as first real measurement, not yet run.	v7 restructure
Descending rounds (plan → architecture → edge cases → code, frozen per level)	■ BUILT	Shipped v7.	Ideas.md
Claim schema w/ typed evidence, bounded tool calls, canary objections	■ BUILT	v7.5 six-item batch.	FOCUS.md
Resource Allocator (route attention to real disagreement)	■ BUILT	v7.3, shipped before v7.5.	FOCUS.md
Ambiguity-union criteria, roster inversion, 2-strong-model comparison chain	■ BUILT	v7.5 batch.	FOCUS.md
Provider failure → recorded dropout, never crash the run	■ BUILT	Shipped as part of v7 fast-follows.	Ideas.md
Unevidenced FAILED verdicts synthesized by our own code	■ BUILT	Root-caused (evidence-driven verdict enforcement) and fixed.	Ideas.md
Partial-output safety check schema-mismatched	■ PLANNED	Named bug — check current stage-contract.js before assuming still open.	v7 backlog #2
Fixed budgets must fail loudly, never truncate silently	■ PLANNED	Recurring defect named twice; verify current budget-handling code.	v7 backlog #3
"Prompts are not validation" systematic enforcement pass	■ PARTIAL	Individual instances fixed; no evidence of a full sweep.	v7 backlog #4
Graphe paranomon challenge stage (reopen a signed-off criterion, at a cost)	■ PLANNED	Ranked #1 mechanism candidate against the debate plateau.	research section

Idea	Status	Notes	Source
Reward information contributed, not agreement (isegoria)	■ PLANNED	Independence-skew instrument exists, not wired to scoring.	research section
Invert speaking order (weakest seats first)	■ IDEA	Cheapest falsifiable experiment, never run.	research section
Sortition (draw seats by lot per run)	■ IDEA	Conflicts with cost control + no-lab-grades-own-draft rule.	research section
Discessio: discrete support/object/abstain as structured data	■ PLANNED	Motivated by two real prose-vs-verdict disagreement bugs.	research section
Per discessionem: skip debate when first pass is unanimous	■ IDEA	Cost-saving idea, not built.	research section
Structured objections, Toulmin-shaped	■ PLANNED	The substrate every later metric needs.	v7 restructure #4
Disagreement-survival / novelty metrics	■ PLANNED	Depends on structured objections above.	v7 restructure #5
"Dissent and what you lose" section in every deliverable	■ PLANNED	Not confirmed built.	v7 restructure #6
Confidence calibration per criterion	■ IDEA	Not scoped.	v7 restructure #7
Cost routing (frontier on debate, cheap on skeleton/build)	■ PLANNED	Aider-style split reports 30-50% cost reduction elsewhere.	v7 restructure #8
MLLM north star: many models propose/vote to build a whole business	■ PLANNED	"Third slice" — gated on tool-verification + descending rounds, both now shipped.	V7 NORTH STAR
Plan → executable, dependency-ordered task list	■ IDEA	Bridges a plan document to real Claude Code work.	v7 additions #1
Follow-up questions against a completed run	■ IDEA	Cheap: retrieval + one seat, answers from the run's own record.	v7 additions #2
Per-run debate-value summary attached to report.json	■ IDEA	Archive-wide script already exists; per-run attach not confirmed.	v7 additions #4
Deliverable in the buyer's language	■ IDEA	Legal texts stay human-written; plan translation is lower risk.	v7 additions #5
Chain blames builder for sections the chain itself requires	■ PLANNED	Cost one revision round on 3 separate runs; cheap fix identified.	Ideas.md
npm publish	■ BUILT	the-high-council@0.7.5 live on registry.	FOCUS.md
Local-model / Ollama seats	■ BUILT	Shipped v7.1.	FOCUS.md
Windows .exe + Linux Appliance, CI workflow	■ PARTIAL	Binaries built + smoke-tested; CI workflow has never actually run.	FOCUS.md
backend-developer / frontend-developer generic skills	■ BUILT	De-SMO'd, MIT, shipped v0.7.0.	FOCUS.md

Idea	Status	Notes	Source
Reusable criteria library (by task shape)	■ IDEA	GP's 40-criterion rubric took weeks; a shared library raises the floor.	<i>Ideas.md</i>
Pre-delivery checklist for customer plans	■ IDEA	Built from known failure modes; not built as a gate.	<i>Ideas.md</i>
Plan-vs-build diff capture (ground truth, zero cost)	■ IDEA	Substitute for a paid eval; not built.	<i>Ideas.md</i>
Intake-form quality as the real lever on plan quality	■ IDEA	Commercial insight, applies directly to the plan-shop intake.	<i>Ideas.md</i>

MLLM Coder (the coding / diff-review pillar)

Idea	Status	Notes	Source
MLLM Coder v1 (debate-gated diff review)	■ BUILT	One diff/file/hunk, no write tool. Planned + built.	<i>FOCUS.md</i>
MLLM Coder v2 (OpenRouter interchangeability)	■ BUILT	Merged 620ae44.	<i>FOCUS.md</i>
MLLM Coder v3 (trust-hazard fact, fallback fix, path-sensitive signoff)	■ BUILT	Items 2/3/4 merged.	<i>FOCUS.md</i>
MLLM Coder v3 items 5+6 (debate-before-propose, security-lens critic)	■ DEFERRED	Would redefine what "critic" means in coder-gate. Reopen trigger: a scoped chain.js stage addition.	<i>coder-gate-v3 deferral</i>
MLLM Coder v4 (preflight, verify re-invocation, --signoff, offline scoring)	■ BUILT	Re-verified from scratch, 40/40 literal acceptance tests.	<i>FOCUS.md</i>
MLLM Coder v4.5 Project B (gate integration vs. stub agent)	■ BUILT	Task 0+1 done; Muad has since added further tasks directly.	<i>FOCUS.md</i>
MLLM Coder v4.5 Project A (real write-capable coding agent)	■ DEFERRED	Deferred, not cancelled. Reopen: ≥3 cited unreviewed-apply failures elsewhere, or a real pilot user demand.	<i>FOCUS.md</i>
MLLM Coder v6 (status ledger, span-tree progress records)	■ BUILT	Items V6-1+V6-2 merged.	<i>git log</i>
MLLM Coder v6 remaining items	■ PLANNED	2/3 signoff, round cap hit. One open objection left as a named human decision.	<i>IDEAS-INDEX</i>
Gesture-based diff review UX	■ IDEA	Belongs to GUI phase 3 — v1 has zero UX surface by design.	<i>FLAGSHIP-IDEAS</i>
Multi-candidate diff comparison (Rover-style)	■ IDEA	Reinvented six times across the OSS ecosystem in 2026, nothing shipped in this shape.	<i>deep-research-v8</i>
Context partitioning across proposal-stage seats	■ PARTIAL	v1 shipped.	<i>git log</i>
Cold-reader coherence check (post-signoff, zero-context)	■ BUILT	Merged.	<i>git log</i>

Idea	Status	Notes	Source
Final security-review gate (read-only, post-build, typed findings)	■ BUILT	Merged.	<i>git log</i>

Packaging, protocol & ecosystem

Idea	Status	Notes	Source
Versioned run/event protocol (events.jsonl, live)	■ IDEA	Cheap, no architecture conflict, useful for Sophi-A's GUI too.	<i>Ideas.md</i>
ACP (Agent Client Protocol) integration	■ IDEA	Real standard; a proxy-chain RFD is "the shape a gate actually is."	<i>deep-research-v8</i>
Reserved reason codes + a named non-verdict for "could not judge"	■ IDEA	Cheapest real item from the v8 followup.	<i>deep-research-v8</i>
Digest-bound one-time approvals for human-stop gates	■ IDEA	From v8 followup.	<i>deep-research-v8</i>
Meter a debate run as a cost unit	■ IDEA	Named as a clean, currently-unfilled gap industry-wide.	<i>deep-research-v8</i>
Monthly research methodology (paper feed, changelog watch)	■ BUILT	Settled, unanimous 3/3.	<i>IDEAS-INDEX</i>
Hash-chained audit-log entries (AgentGuard-style)	■ IDEA	Top find from market survey, unbuilt.	<i>market-feature-survey</i>
Academic citations for the blind-panel approach	■ BUILT	Built into v5 item 6.	<i>IDEAS-INDEX</i>
Capability-first framing for policy checks	■ BUILT	Built into v5 item 5.	<i>IDEAS-INDEX</i>

Strategy & positioning

Idea	Status	Notes	Source
Durable moat is not the algorithm (MIT + commodity models)	■ IDEA	Possibly the single most consequential open question the company has.	<i>IDEAS-INDEX</i>
Candidate moat: orchestration UX (Sophi-A) as a craft moat	■ IDEA	Positioning idea.	<i>IDEAS-INDEX</i>
Candidate moat: honestly-published trust record over time	■ PARTIAL	Discipline already practiced; SWE-bench pilot is the first visible step.	<i>IDEAS-INDEX</i>
Candidate moat: curatorial taste made visible (public "why we said no")	■ IDEA	Nobody in the space publishes declined-idea reasoning as a selling point.	<i>IDEAS-INDEX</i>

Idea	Status	Notes	Source
"Beats a single model" requires 3 conditions (decorrelation, aggregator, routing)	■ BUILT	Promoted as a standing positioning check before any marketing claim.	<i>mllm-beyond-single-model-ideas</i>
Sophi-A/THCMCP merge into one runtime+clients product	■ DECLINED	Directly contradicts THCMCP's own CLAUDE.md — "do not conflate them."	<i>Ideas.md</i>
Trust/risk-tier layer on agent actions	■ PLANNED	Not new — a second description of the open required_signoff_paths gap.	<i>FOCUS.md item 6</i>
Subscription-seat-as-third-party-automation	■ DECLINED	Against Anthropic's own ToS, confirmed via a closed GitHub issue.	<i>subscription-seat-to-s-check.md</i>

Generated from THCMCP/IDEA-MATRIX.md, 2026-09-16. Not a plan or commitment — status view of ideas gathered across Review/ research docs, relay runs, and chat. See the source .md for full pointers and reasoning per row.